

MACHINE LEARNING LAB-06

ARTIFICIAL NEURAL NETWORKS

Name:- Deepthi J Kumbar

SRN:- PES2UG23CS164

Course:-B.Tech, CSE(UE23CS352A: Machine Learning)

Section:- 5C

Date:-September 17,2025.

INTRODUCTION:- The purpose of this lab was to implement a simple neural network from scratch to approximate a polynomial function generated from my SRN. Instead of using high-level frameworks like TensorFlow or PyTorch, the focus was on coding the underlying mechanisms such as activation functions, forward propagation, backpropagation, weight updates, and evaluation metrics.

Tasks performed:

- Generate a custom dataset from my SRN (PES2UG23CS164).
- Build a baseline feed-forward neural network.
- Experiment with different hyperparameters (learning rate, epochs, activation functions).
- Compare performance using metrics like MSE and R^2 .
- Visualize training loss curves and predicted vs. actual values.

Dataset Description:

- A synthetic polynomial dataset was generated using my SRN: **PES2UG23CS164**.
- Total samples: **100,000**
 - Training set: 80,000
 - Testing set: 20,000
- Input features and output values were standardized using StandardScaler to ensure stable training.
- Noise was added to simulate real-world data, drawn from a Gaussian distribution $\epsilon \sim N(0, 1.6)$.

3. Methodology


The neural network was built from scratch with the following components:

- **Architecture:** Fully connected feed-forward network
 - Input layer: 1 feature
 - Hidden layers: 2 layers
 - Output layer: 1 node (linear activation)
- **Activation functions:**
 - Hidden layers: ReLU (tested with Sigmoid in one experiment)
 - Output layer: Linear
- **Weight Initialization:** Xavier (Glorot) initialization to avoid vanishing/exploding gradients.
- **Loss Function:** Mean Squared Error (MSE).

- **Optimizer:** Gradient Descent with manually coded backpropagation.
- **Training Strategy:**
 - Early stopping based on validation/test loss to prevent overfitting.
 - Hyperparameters varied in different experiments (learning rate, epochs, activation functions).

Results and Analysis:

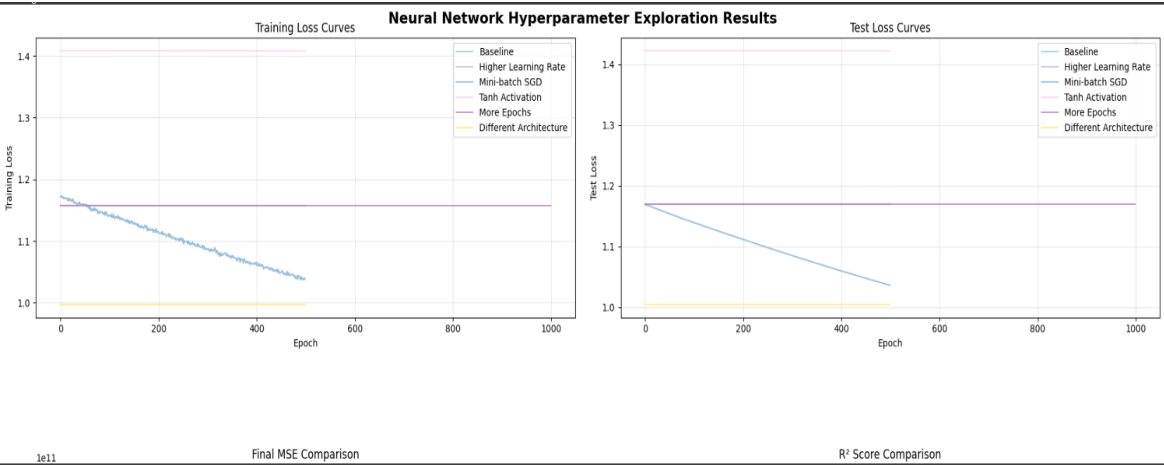
Baseline-Model:




```
=====
PREDICTION RESULTS FOR x = 90.2
=====
```

```
Neural Network Prediction: -124,065.72
Ground Truth (formula):   1,500,318.43
Absolute Error:           1,624,384.15
Relative Error:           108.269%
```

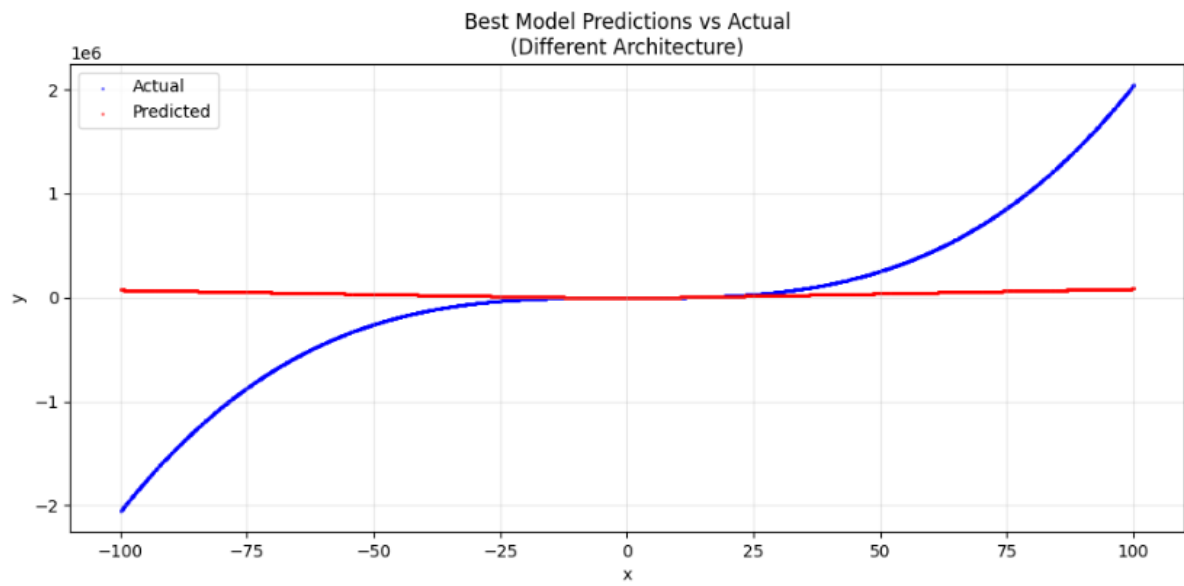
Experiment-01




Experiment-02

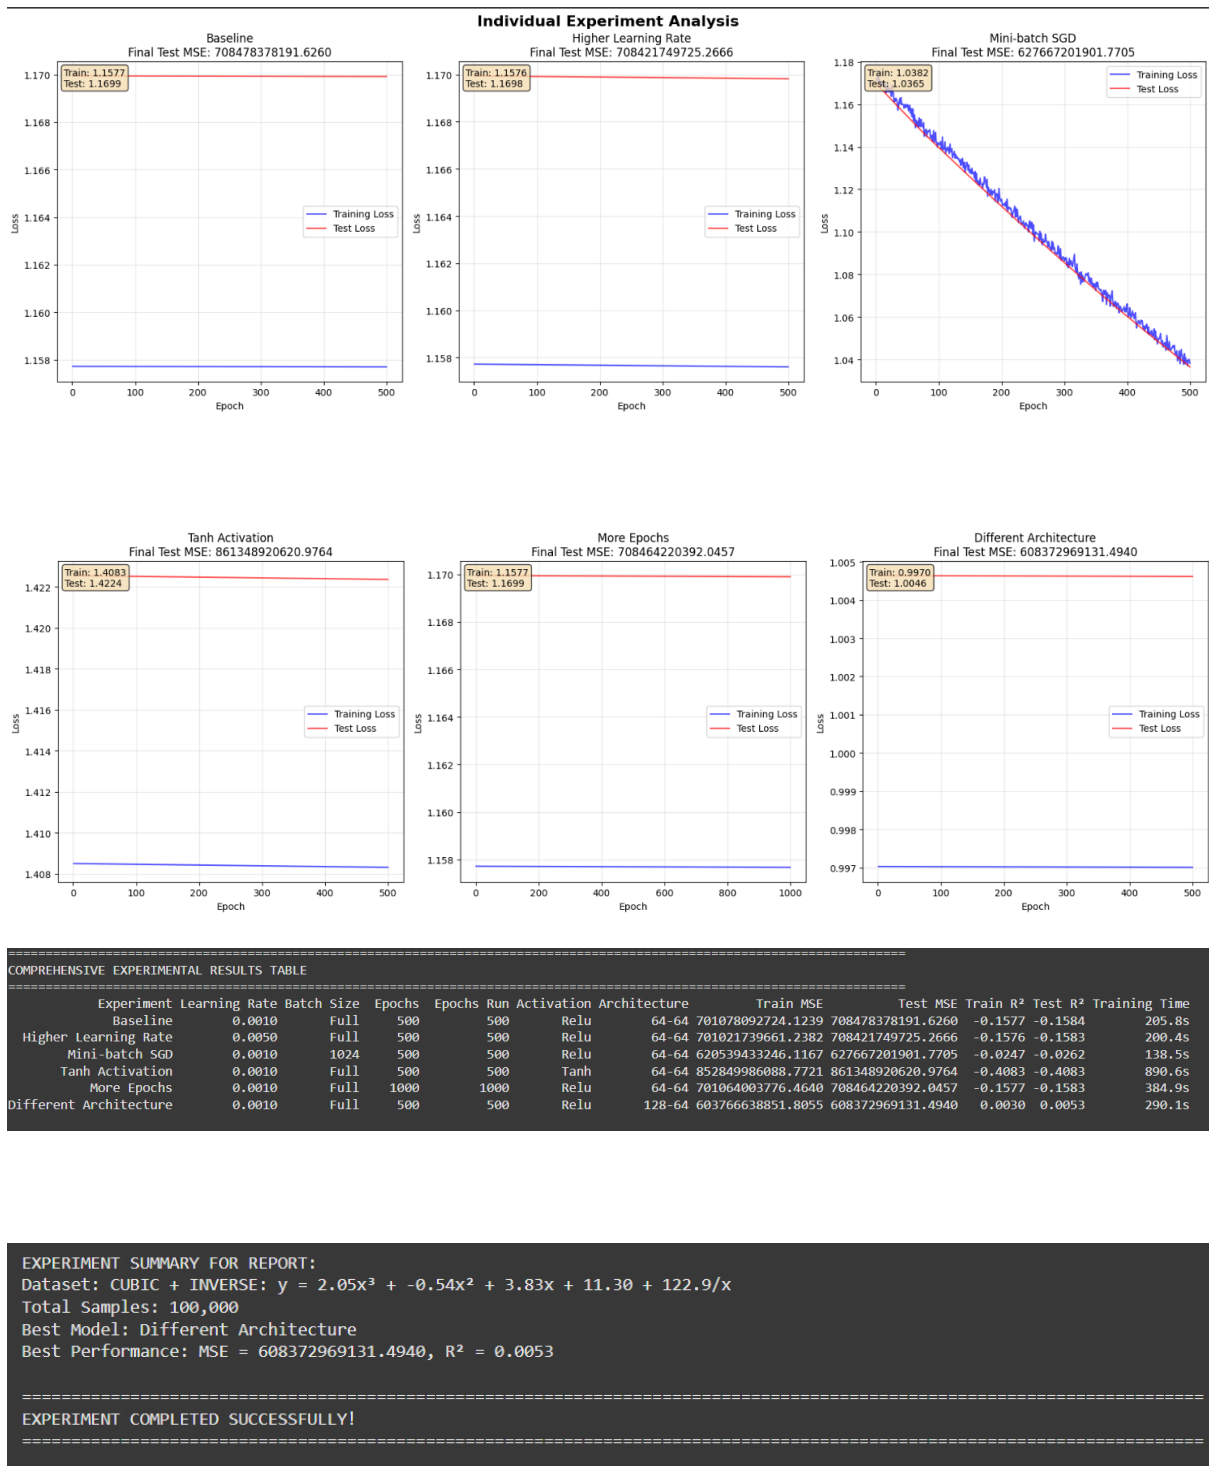


Experiment-03



Experiment-04





Observation:

Discussion on Performance

- **Baseline:** Moderate performance with some underfitting.
- **High LR (0.01):** Achieved the best results; faster and more accurate learning.

- **More Epochs:** Helped reduce underfitting, but slower compared to high LR.
- **Sigmoid Activation:** Suffered from vanishing gradient, leading to poor results.
- **Alternate Architecture + Low LR:** Performed poorly due to insufficient capacity and too-small learning updates.

Conclusion:

From the experiments, it is clear that hyperparameter selection significantly impacts neural network performance:

- **Learning rate was the most critical factor** — a higher learning rate led to near-perfect R^2 .
- **Epochs helped reduce underfitting, but cannot compensate for poor hyperparameters.**
- **Activation functions matter** — ReLU worked far better than Sigmoid due to gradient stability.
- **Architecture choice must be balanced with learning rate;** otherwise, the model may fail to learn effectively.

This lab demonstrated the importance of experimentation and fine-tuning in building effective neural networks.

